

Neuronized Priors Project

Christopher Crabtree, Lanxin Yang, Yunshan Duan

December 14, 2021

Abstract

This paper presents a description and replication of the primary results of “Neuronized Priors for Bayesian Sparse Linear Regression” by Minsuk Shin and Jun S. Liu. We implemented the MCMC sampling algorithm and the MAP optimization algorithm for the neuronized priors proposed in this paper. And we also implemented the methods of their standard shrinkage prior counterparts, including the SpSL, Bayesian LASSO, and horseshoe priors. Simulation studies and real data examples are investigated, and comparison between optimization methods are conducted. The results are parallel to those in the original paper, that neuronized priors yields similar performance as the standard shrinkage priors, and have computational advantage by reducing the complexity when sampling from the posteriors.

1 Introduction

Consider the standard linear regression model

$$\mathbf{y} = X\boldsymbol{\theta} + \boldsymbol{\epsilon},$$

where $\boldsymbol{\epsilon} \sim N(0, \sigma^2 I)$. $\mathbf{y} = (y_1, \dots, y_n)^T$ is the vector of responses, X is the $n \times p$ covariate matrix, $\boldsymbol{\theta} = (\theta_1, \dots, \theta_p)^T \in \mathbb{R}^p$ is the coefficient vector.

To assume sparsity on the prior of $\boldsymbol{\theta}$, a shrinkage prior is often imposed on θ_j 's. One popular class of shrinkage prior is the one-group continuous shrinkage prior, which is a hierarchical scale-mixture of Gaussian distributions.

$$\theta_j | \nu_w^2, \tau_j^2 \sim N(0, \nu_w^2 \tau_j^2),$$

where $\tau_j^2 \sim \pi_\tau$ and $\nu_w \sim \pi_g$ are pre-specified. A few choices of the distributions have been shown to induce desirable shrinkage, including the Strawderman-Berger prior, the **Bayesian Lasso** [Park and Casella, 2008], the **horseshoe** prior [Carvalho et al., 2010], the generalized double Pareto, and the Dirichlet-Laplace prior.

Another popular class of shrinkage prior is the **spike-and-slab (SpSL)** priors [Mitchell and Beauchamp, 1988, George and McCulloch, 1993], also known as two-group mixture

priors. It can be written as a mixture of the “spike” distribution π_0 , highly concentrated around zero, and the “slab” distribution π_1 , which is relatively disperse.

The paper [Shin and Liu, 2021] proposes **neuronized priors**, in which each regression coefficient is reparameterized as a product of a weight parameter and a transformed scale parameter via an activation function. Popular shrinkage priors fit this unified form of neuronized priors, including SpSL, Bayesian Lasso, Cauchy, horseshoe, ect.

It is demonstrated in Shin and Liu [2021] that the variable selection procedures based on the neuronized priors have the advantage both computationally and theoretically. The characteristic of unifying various classes of shrinkage priors allow one to test out different priors simply by changing the activation function. Additionally, posterior inference of neuronized priors does not rely on the prior-likelihood conjugacy, leading to comparable or better efficiency in computation compared to the standard priors. Moreover, with the conditions put on the activation function and the hyperpriors, the neuronized Bayesian regression achieves the optimal posterior contraction rate and the RWMH algorithm converges to the target distribution at an exponential rate.

The rest part of the paper is organized as follows. Section 2 gives the detailed definition of the neuronized priors. Section 3 shows a few popular shrinkage priors including the SpSL, Bayesian Lasso, horseshoe, and Cauchy priors, and their neuronized counterparts. Section 4 illustrates the computational strategies of posterior inference and the advantages and the neuronized priors. Simulation studies are conducted in Section 5 to compare the results of inference based on different priors and their neuronized counterparts. Two real data examples are also analyzed in Section 5. And conclusions are given in Section 6.

2 Neuronized prior

Definition 2.1 (Neuronized prior). A neuronized prior is for θ_j is defined as

$$\theta_j := T(\alpha_j - \alpha_0)w_j, \quad (1)$$

where T is a non-decreasing activation function, the scale parameter α_j follows $N(0, 1)$, and the weight parameter w_j follows $N(0, \tau_w^2)$, all independently for $j = 1, \dots, p$.

Under the neuronized prior for $\boldsymbol{\theta}$, and the prior $\pi(\sigma^2)$ for σ^2 , the joint posterior distribution can be written as

$$\pi(\boldsymbol{\alpha}, \mathbf{w}, \sigma^2 | \mathbf{y}, \alpha_0) \propto \frac{\pi(\sigma^2)}{\sigma^n} \exp\left\{-\frac{\|\mathbf{y} - X\boldsymbol{\theta}(\boldsymbol{\alpha}, \mathbf{w}, \alpha_0)\|_2^2}{2\sigma^2} - \frac{\boldsymbol{\alpha}^T \boldsymbol{\alpha}}{2} - \frac{\mathbf{w}^T \mathbf{w}}{2\tau_w^2}\right\}, \quad (2)$$

where $\boldsymbol{\alpha} = \{\alpha_1, \dots, \alpha_p\}^T$, $\mathbf{w} = \{w_1, \dots, w_p\}^T$, and

$$\boldsymbol{\theta}(\boldsymbol{\alpha}, \mathbf{w}, \alpha_0) = \{T(\alpha_1 - \alpha_0)w_1, \dots, T(\alpha_p - \alpha_0)w_p\}^T \triangleq D_{\boldsymbol{\alpha}} \mathbf{w},$$

in which $D_{\boldsymbol{\alpha}}$ is the diagonal matrix with diagonal elements the $T(\alpha_j - \alpha_0)$'s.

3 Neuronization of Standard Sparse Priors

3.1 Discrete and Continuous SpSL Priors

The spike-and-slab priors [Mitchell and Beauchamp, 1988, George and McCulloch, 1993] are defined as

$$\begin{aligned}\theta_j|\gamma_j &\sim (1 - \gamma_j)\pi_0(\theta_j) + \gamma_j\pi_1(\theta_j) \\ \gamma_j &\sim \text{Bernoulli}(\eta)\end{aligned}\tag{3}$$

When π_0 is the point-mass at zero, we call the prior a discrete SpSL prior, otherwise we call it a continuous prior. Typically, π_0 and π_1 for a continuous SpSL prior can be defined as two Gaussian distributions with a small and a large variance, respectively.

The computation of the posterior inference based on the discrete SpSL prior is often challenging. Several MCMC sampling strategies and stochastic search strategies have been proposed in literature to counter the difficulty. The MCMC sampling of the posterior using the continuous prior can be more efficiently implemented because of the advantage of conjugacy. However, compared to discrete SpSL prior, the continuous prior does not automatically conclude selected coefficients, and extra variable selection steps are required.

Let the activation function in Equation (1) be the rectifier linear unit (ReLU) function,

$$T(t) = \max\{0, t\}.$$

Set $\alpha_0 = 0$, then the activation function $T(\alpha_j - \alpha_0)$ follows an equal mixture of a point mass at zero and the half standard Gaussian. Under this setting, the neuronized prior in Equation (1) is equivalent to the discrete SpSL distribution with $\eta = \frac{1}{2}$, which is in the form

$$\begin{aligned}\theta|\gamma &\sim (1 - \gamma)\delta_0(\theta) + \gamma\pi(\theta), \\ \gamma &\sim \text{Bernoulli}\left(\frac{1}{2}\right),\end{aligned}\tag{4}$$

Figure 1 shows the histogram of $T(\alpha)$, $T(\alpha)w$, and the standard SpSL prior. Here $\tau_w^2 = 1$. As is shown in the figure, the neuronized prior using the ReLU function as activation function is equivalent to the discrete SpSL prior.

Generally, without the assumption $\alpha_0 = 0$, the neuronized prior with the ReLU activation function gives the same distribution as the standard SpSL prior in Equation (1) when γ is set to be $\gamma \sim \text{Bernoulli}(\Phi(-\alpha_0))$, with Φ denoting the standard Gaussian CDF. The hyper-parameter α_0 controls the probability of sparsity of the prior distribution,

$$P(T(\alpha_j - \alpha_0) = 0|\alpha_0) = P(\alpha_j < \alpha_0|\alpha_0) = \Phi(\alpha_0).$$

Conversely, it means that we can choose a desirable level of sparsity with $\alpha_0 = -\Phi^{-1}(\eta)$, for any $\eta \in (0, 1)$.




Figure 1: Histogram of $T(\alpha)$, $T(\alpha)w$, and the standard SpSL prior.

For standard discrete SpSL distribution, as is shown in Scott and Berger [2010], a Beta hyper-prior can be put on the sparsity parameter η ,

$$\eta \sim \text{Beta}(a_0, b_0), \quad (5)$$

where $a_0 = b_0 = 1$ leading to a strong effect on multiplicity correction. If $(a_0, b_0) = (1, p^a)$ for $a > 1$, and the number of predictors p increases at a sub-exponential rate of n , $p = O(\exp(n^c))$ for $c < 1$, Castillo and van der Vaart [2012] and Castillo et al. [2015] propose that this SpSL procedure achieves model selection consistency and the optimal posterior contraction rate. By adopting a hyper-prior on α_0 ,

$$\pi(\alpha_0) \propto \Phi(-\alpha_0)^{a_0-1} (1 - \Phi(-\alpha_0))^{b_0-1} \phi(\alpha_0),$$

the neuronized SpSL prior distribution is identical to the standard SpSL prior in Equation (1) with the Beta prior (5) on η .

Using the “leaky” ReLU activation function

$$T(t) = \max\{ct, t\},$$

where c is a constant $c < 1$, we can obtain the continuous SpSL prior from Equation (1).

3.2 Bayesian Lasso

The Bayesian Lasso Park and Casella [2008] extends the original Lasso model by Tibshirani et al. Tibshirani [1996]. by using a Laplace prior on the coefficients β . In order to efficiently compute this, the Park et al. use a decomposed representation of the Laplace distribution as follows:

$$\frac{a}{2}e^{-a|\beta_i|} = \int_0^\infty \left(\frac{a^2}{2}e^{-a^2s/2}\right) \left(\frac{1}{\sqrt{2\pi s}}e^{-z^2/(2s)}\right) ds \quad (6)$$

This representations allows for each β_i to be given the following distribution:

$$\begin{aligned} \beta_i &\sim \mathcal{N}(0, \sigma^2 \tau_i) \\ \sigma^2 &\sim \pi(\sigma^2) d\sigma^2 \\ \tau_i &\sim \frac{\lambda^2}{2} e^{-\lambda^2 \tau_i^2 / 2} d\tau_i \end{aligned}$$

With this representation, the shrinkage of model coefficients can be controlled with 2 parameters: a global shrinkage parameter λ^2 , and a parameter τ_i^2 that is local to each coefficient. This representation is computationally convenient as the conditional distributions of β and σ^2 and τ_i^2 have simple multivariate normal, inverse-gamma, and inverse-gaussian distributions respectively. Likewise updating of the penalty parameter λ can be done in a Bayesian manner, but in this paper the authors opted to use the estimate obtained from standard Lasso cross-validation.

While this representation of the coefficient prior allows for a computationally feasible method for obtaining posterior estimates, optimizations cannot translate to other sparse prior distributions. For example, improvements such as inference with a reversible jump MCMC algorithm Chen et al. [2011] do not directly apply to other sparse priors. With this in mind, Shin et al. present a neuronized version of the Bayesian Lasso prior.

Neuronization can be achieved simply by setting the activation function to: $T(t) = t$. This is motivated by previous work such as Hoff et al. Hoff [2017] which noted that the product representation of the prior on β_i is similar to the Bayesian lasso.

3.3 Horseshoe, Cauchy and Their Generalization

As a continuous shrinkage prior, horseshoe prior one of the most popular local-global prior to reach the purpose of locally adaptive shrinkage because of its sharp asymptotic at the origin as well as its heavy tail. The density function horseshoe prior can be presented as a mixture of Gaussian distributions:

$$\theta_i | \nu_w^2, \tau_i^2 \sim \mathcal{N}(0, \nu_w^2 \tau_i^2) \quad (7)$$

$$\tau_i \sim \pi_\tau \equiv C^+(0, 1) \quad (8)$$

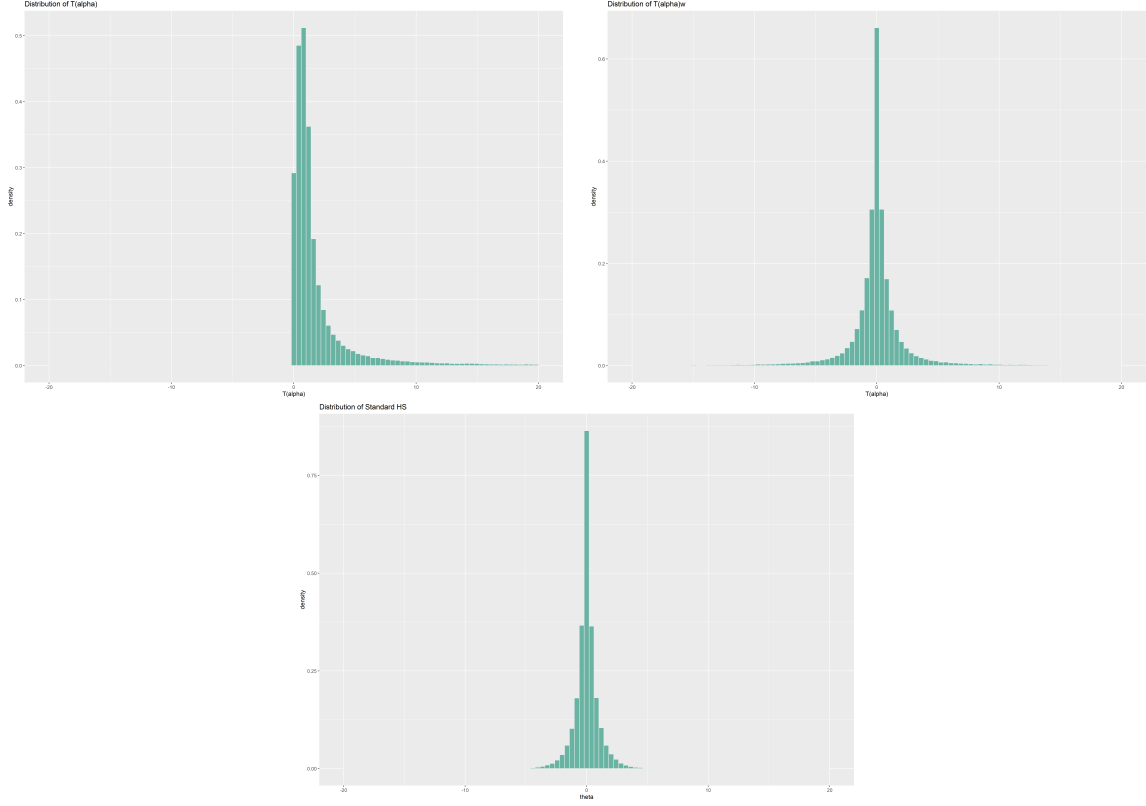


Figure 2: Histogram of $T(\alpha)$, $T(\alpha)w$, and the standard Horseshoe prior.

π_τ here is a half-Cauchy distribution for standard deviation τ . To fit horseshoe priors, we choose to construct heavy tails of the local shrinkage prior by transform a Normal distribution to a heavy tail distribution through an exponential function, and then calculate activation function that can make the transformed priors approximate the horseshoe prior.

$$T(t) = \exp(\lambda_1 \text{sign}(t)t^2), \lambda \in (0, 1) \quad (9)$$

For $U = T(Z)$ and $Z \sim N(0, 1)$, then the density function of U would be:

$$f_U(u) \propto u^{-1 - \frac{\text{sign}(\log(u))}{2}} |\log(u)|^{\frac{1}{2}} \quad (10)$$

This exponential activation function $T(t)$ construct a polynomial -tailed prior on $\theta \sim T(\alpha)\omega$, which align with our purpose. Figure 2 show the density of $T(\alpha)\omega$ along with θ , which is almost same with the original horseshoe prior distribution.

In order to generalize horseshoe priors, as λ_1 indicates the tail behavior of the current neuronized prior, we add two more polynomial terms with lower exponents to the

exponential activation function:

$$T(t) = \exp(\lambda_1 \text{sign}(t)t^2 + \lambda_2 t + \lambda_3) \quad (11)$$

By numerical calculation, when $T(t) = \exp(0.5 \text{sign}(t)t^2 + 0.733t)$, the neuronized prior approximate horseshoe prior well and $T(t) = \exp(0.5t^2 - 1.27t + 0.29)$ approximate to standard Cauchy priors Shin and Liu [2021]. However, the generalized activation function for Cauchy will change to $T(t) = \exp(\lambda_1 t^2 + \lambda_2 t + \lambda_3)$ since this activation function will have stronger shrinkage effects for weak signals.

Under this circumstance, the activation function solve the challenge that the corresponding MCMC algorithm can't achieve to geometric ergodicity under horseshoe prior, which is a heavy tail prior Mengersen and Tweedie [1996]. The MCMC procedure can achieve convergence more efficiently with the activation function, because the tail is stable in this case.

4 Managing Neuronized Priors

4.1 Find the Activation Function to Match a Given Prior

In order to find an activation function $T(t)$ for the neuronized prior that discussed in section 3, the author introduce a family of $T(t)$, T_ϕ , and then find the ϕ which could minimize the distance between the result neuronized prior and the target distribution $\pi(\theta)$. To enhance flexibility, the function space spanned by a class of B-spline basis function $B(t)$, B is a vector of K dimension rational space where $\phi \in R^K$. Consequently, $T_\phi(t) = B(t)\phi$.

The essay use neuronized horseshoe prior as an example, whose activation function is $T(t) = \exp(\lambda_1 \text{sign}(t)t^2 + \lambda_2 t + \lambda_3)$. In this case, $\phi = \lambda_1, \lambda_2, \lambda_3$. The class of activation function for horseshoe now become:

$$T_\zeta(t) = \exp(\lambda_1 \text{sign}(t)t^2 + \lambda_2 t + \lambda_3) + B(t)\phi \quad (12)$$

$$\zeta = \{\lambda_1, \phi\} \quad (13)$$

As the determinant parameter of prior's tail behavior, λ_1 can be found to match the tail of the horseshoe prior base on the inequality property: For $\kappa > 0$, c_1, c_2 are some positive constants, there exist:

$$c_1(\log|\theta|)^{-\frac{1}{2}}|\theta|^{(-1-\frac{1}{2\lambda_1})(1+\kappa)} \leq \pi(\theta) \leq c_2(\log|\theta|)^{-\frac{1}{2}}|\theta|^{(-1-\frac{1}{2\lambda_1})(1+\kappa)} \quad (14)$$

After λ_1 is fixes, we generate a large number (S) of samples from $\tilde{\theta}_{i\zeta,i} = T_{\zeta,\phi}(\alpha_i - \alpha_0)\omega_i$ and $\theta_i \sim \pi(\theta)$, $i = 1, 2, \dots, S$. Then by using a search algorithm to find ζ to minimize $\sum_{i=1}^S |\tilde{\theta}_{\zeta}^{(i)} - \theta^{(i)}|$ Kirkpatrick et al. [1983], where $\theta_{\zeta}^{(i)}, \theta^{(i)}$ are the i th largest number of the generated samples.

4.2 Choosing Hyper-Parameters

There are two hyper-parameter that need to be determined: the global shrinkage parameter τ_w^2 and the bias parameter α_0 . As for the value of α_0 , there are several situation:

- For continuous neuronized prior, α_0 is set to the default value 0.
- For discrete Spsl prior, where $P(T(\alpha_j - \alpha_0) = 0 | \alpha_0) = \Phi(\alpha_0)$. As the author stated for the discrete Spsl prior, a hyper-prior can be imposed on α_0 so the the sparsity level is determined by the data set.
- Sampling α_0 on other parameter in Gibbs sampling and RWMH algorithm is highly inefficient due to the high correlation of its posterior with that of other parameters.

When choosing value for τ_w^2 :

- When $E(T^2(\alpha))$ is bounded, the value of τ_w^2 need to reflect the signal strength in the data. However, a theoretical justified selection for τ_w^2 has not been found.
- When $E(T^2(\alpha))$ doesn't bound such as horseshoe and Cauchy, we first find a shrinkage factor κ_j : $E_j(\kappa_0) = \min(0.01, 0.1 \times n/p)$, which determine the shrinkage level of θ_j . Then we calculate τ_w^2 by

$$\tau_w^2 = \frac{1 - \kappa_j}{\kappa T^2(\alpha_j)}$$

5 Sampling and Optimization With Neuronized Priors

5.1 MCMC Sampling with Neuronized Priors

Under the linear regression model mentioned previously, the joint posterior distribution of the parameters follows the form of Equation (2). Therefore, the full conditions can be written as follows. The conditional posterior distribution of \mathbf{w} . given other parameters is a Gaussian distribution,

$$\mathbf{w} | \mathbf{y}, \boldsymbol{\alpha}, \sigma^2, \tau_w^2 \sim N(\tilde{\boldsymbol{\mu}}, \sigma^2 \tilde{\boldsymbol{\Sigma}}), \quad (15)$$

where $\tilde{\boldsymbol{\Sigma}} = (D_\alpha X^T X D_\alpha + \sigma^2 \tau_w^{-2} I)^{-1}$, and $\tilde{\boldsymbol{\mu}} = \tilde{\boldsymbol{\Sigma}} D_\alpha X^T \mathbf{y}$, and D_α is defined as in Section 2. However, when the number of coefficients p is large relative to the number of observations n , the computational complexity of solving $(D_\alpha X^T X D_\alpha + \sigma^2 \tau_w^{-2} I)^{-1}$ is highly expensive. Shin and Liu [2021] adopts the fast sampling procedure proposed in Bhattacharya et al. [2016], which reduces the complexity from $O(p^3)$ to $O(n^2 p)$.

Because w_j and α_j are highly correlated a posteriori, the algorithm draw samples from the posterior through the following scheme. For $j = 1, \dots, p$, we draw sample α_j conditional on $\boldsymbol{\alpha}_{-j}$ and \mathbf{w} , that is, $\alpha_j^* \sim \pi(\alpha_j | \mathbf{y}, \mathbf{w}_{-j}, \boldsymbol{\alpha}_{-j})$ by a RWMH algorithm, based on the log-target function,

$$-\log(v_j)/2 - \alpha_j^2/2 + v_j m_j^2/(2\sigma^2), \quad (16)$$

where $v_j = X_j^T X_j T(\alpha_j - \alpha_0) + \sigma^2 / \tau_w^2$, and $m_j = \mathbf{r}^T X_j T(\alpha_j - \alpha_0) / v_j$. And then sample w_j from $\pi(w_j | \mathbf{y}, \boldsymbol{\alpha}_j, \alpha_j^*)$, conditional on other parameters.

$$w_j | \mathbf{y}, \boldsymbol{\alpha}_{-j}, \alpha_j, \mathbf{w}_{-j}, \sigma^2, \tau_w^2 \sim N(m_j, \sigma^2 v_j^{-1}) \quad (17)$$

For continuous neuronized priors, α_0 is set to be 0 and for neuronized SpSL, a hyper parameter is imposed on α_0 which is discussed in the next section. We implemented this algorithm with the help of the NPrior package included in Shin and Liu [2021], and conducted simulation and real data examples in later sections. The detailed algorithm is shown in Algorithm 1.

Algorithm 1 A general MCMC algorithm for neuronized priors

Require: Initialization of $\boldsymbol{\alpha}, \alpha_0, \mathbf{w}, \sigma^2$.

Sample \mathbf{w} conditional on $\mathbf{y}, \boldsymbol{\alpha}, \sigma^2$ from Equation (15).

for $i = 1, \dots, N$ **do**

Set $\mathbf{r} = \mathbf{y} - X\boldsymbol{\theta}(\boldsymbol{\alpha}, \mathbf{w})$.

for $j = 1, \dots, p$ **do**

Update $\mathbf{r} = \mathbf{r} + X_j T(\alpha_j - \alpha_0) w_j$.

for $irep = 1, \dots, N$ **do**

Sample α_j from $\alpha_j | \mathbf{y}, \boldsymbol{\alpha}_{-j}, \mathbf{w}_{-j}, \sigma^2, \tau_w^2$ using a RWMH step with Equation (16).

Sample w_j from $w_j | \mathbf{y}, \boldsymbol{\alpha}_{-j}, \alpha_j, \mathbf{w}_{-j}, \sigma^2, \tau_w^2$ with Equation (17).

end for

Update $\mathbf{r} = \mathbf{r} - X_j T(\alpha_j - \alpha_0) w_j$.

end for

Sample σ^2 from $\sigma^2 | \mathbf{y}, \boldsymbol{\alpha}, \mathbf{w}, \tau_w^2$.

Sample δ based on the method described in Section 5.2.

end for

5.2 Sample α_0 efficiently

Because of the high correlation between α_0 and the α_j 's, a naive MH approach in which α_0 is updated by a MH step conditioned on $\boldsymbol{\alpha}$ is highly inefficient. Therefore, Shin and Liu [2021] adopts a group-move via the generalized Gibbs sampling formulation proposed in Liu and Sabatti [2000] which updates $\boldsymbol{\alpha}$ and α_0 simultaneously by a common shift $\delta \in \mathbb{R}$. Specifically, $(\boldsymbol{\alpha}, \alpha_0)$ is updated in the way that

$$(\boldsymbol{\alpha}, \alpha_0) \rightarrow (\boldsymbol{\alpha} + \delta \mathbf{1}, \alpha_0 + \delta),$$

where δ is sampled from $g(\delta) \propto \pi^*(\boldsymbol{\alpha} + \delta \mathbf{1}, \alpha_0 + \delta)$, and $\pi^*(\boldsymbol{\alpha}, \alpha_0)$ is the conditional posterior density for $\boldsymbol{\alpha}$ and α_0 .

When $a_0 = b_0 = 1$ in the prior, $g(\delta) = N((\boldsymbol{\alpha}^T \boldsymbol{\alpha} + \alpha_0)/(p+1), (p+1)^{-1})$. Otherwise, an extra approximation step is needed. The algorithm in [Shin and Liu, 2021] uses a multiple-trial MH independence sampler (MTM-IS) to sample δ . It proposes multiple candidates δ_i 's drawn independently from a proposal distribution and then chooses one from them with probability proportional to their importance weights.

5.3 Computational advantage of neuronized discrete SpSL

The paper [Shin and Liu, 2021] shows that a fully collapsed and a half collapsed Gibbs sampler can be applied to efficiently sample from the posterior with the standard discrete SpSL priors. However, both of the methods are not feasible when the continuous parameters cannot be analytically integrate out. It demonstrates that the neuronized priors are computationally efficient while achieving the same effect as the standard discrete SpSL priors, even if the marginals cannot be derived analytically. As is shown in the proposition in the paper, with a ReLU activation function in the neuronized prior, the conditional distribution $\pi(\alpha_j | \mathbf{y}, \mathbf{w}_{(-j)}, \boldsymbol{\alpha}_{(-j)})$ is a mixture of two truncated Gaussians and therefore can be sampled exactly.

Another advantage of using the ReLU activation function is that when sampling \mathbf{w} , the conditional posterior distribution can be decomposed as a product of independent Gaussian densities so that the numerical inversion of matrix to compute $\tilde{\Sigma}$ in Equation (15) can be avoided.

5.4 A Scalable Algorithm for Finding Posterior Modes

For large dataset, the computation of MCMC algorithm may not be practical so that we may consider optimization-based algorithms. Several optimization methods have been proposed to conduct the posterior inference with the SpSL priors. For example, EMVS [Ročková and George, 2014] and SSLASSO [Ročková and George, 2018] are two optimization procedures for the Bayesian SpSL variable selection problem.

Shin and Liu [2021] proposes a coordinate-ascent algorithm for neuronized (CAAN) priors which finds the MAP estimator. To reduce the chance of being trapped in a local optimum, CAAN uses a warm start strategy by initializing the hyper-parameter leading to a weak shrinkage and increasing the strength of shrinkage gradually.

The algorithm uses a temperature scheme with temperatures t_k such that $t_0 \leq \dots \leq t_{2L+1}$. At each temperature t_k , the coordinate ascent iterations are conducted M times. The CAAN algorithm updates the value of α_j by optimizing it when fixing other parameters $\boldsymbol{\alpha}_{-j}$ and \mathbf{w} . This optimization is easily feasible because the log posterior with respect to α_j is a simple linear combination of a quadratic function and a function of $T(\alpha_j - \alpha_0)$. And \mathbf{w} is updated taking the advantage of conjugacy. To get rid of trapping in a local optima, a random noise $\xi \sim \exp(1)$ is added in the first L levels of temperature scheme. By default, we set $M = 20$, $L = 10$, $N = 20$, $t_k = (3 - 2k/L)^2$ for $k = 1, \dots, L$, and

$t_{L+1} = \dots = t_{2L}$. Detailed algorithm is provided in Algorithm 2 . We implemented this algorithm and conducted simulation studies using it in later sections.

Algorithm 2 The coordinate-ascent algorithm for neuronized prior (CAAN)

Require: Initialization of $\boldsymbol{\alpha}, \alpha_0, \boldsymbol{w}, \sigma^2$.
Set a candidate set of temperature, $\{t_{(1)}, \dots, t_{(2L+1)}\}$.
for $l = 1, \dots, 2L + 1$ **do**
 Set $t = t_{(l)}$.
 Set $\boldsymbol{r} = \boldsymbol{y} - X\boldsymbol{\theta}(\boldsymbol{\alpha}, \boldsymbol{w})$.
 for $irep = 1, \dots, M$ **do**
 for $j = 1, \dots, p$ **do**
 Update $\boldsymbol{r} = \boldsymbol{r} + X_j T(\alpha_j - \alpha_0) w_j$.
 Update α_j by optimizing the log marginalized posterior in Equation (16) with respect to α_j .
 Update w_j from $w_j | \boldsymbol{y}, \boldsymbol{\alpha}_{-j}, \alpha_j, \boldsymbol{w}_{-j}, \sigma^2, \tau_w^2$ with Equation (17).
 Update $\boldsymbol{r} = \boldsymbol{r} - X_j T(\alpha_j - \alpha_0) w_j$.
 end for
 Every N iterations,
 Update $\sigma^2 = (\|\boldsymbol{y} - X\boldsymbol{\theta}(\boldsymbol{\alpha}, \boldsymbol{w})\|_2^2 / t + 2b_1) / \{n + 2a_1 + 2\}$.
 if $l \leq L$ **then**
 Set $\sigma^2 = \sigma^2 + \xi$, where $\xi \sim \exp(1)$.
 end if
 Update $\alpha_0 = \{\sum_{j=1}^p 1(\alpha_j > \alpha_0) + a_0 - 1\} / (p + b_0 - 2)$.
 end for
end for

5.5 Comparison With Other Posterior Optimization Procedures

We conduct simulation to compare the performance of the proposed CAAN and other optimization methods finding the MAP, parallel to what the paper did. We evaluate the property of the CAAN algorithm [Shin and Liu, 2021] and the EMVS [Ročková and George, 2014], which evaluates the MAP estimator Based on an EM formulation. Two metrics are investigated in the study, the mean-squared error (MSE) and the extended Bayesian information criterion (EBIC), which is,

$$\text{EBIC} = \text{BIC} + \zeta |\boldsymbol{k}| \log p,$$

where \boldsymbol{k} is the set of selected variables, ζ is a tuning parameter (here, $\zeta = 1$), and BIC is the Bayesian information criterion.

We generate synthetic data by setting the error variance $\sigma^2 = 1$, and the first 10 elements in the coefficient vector to be randomly 2 or -2, and the rest 0. For EBIC, we use

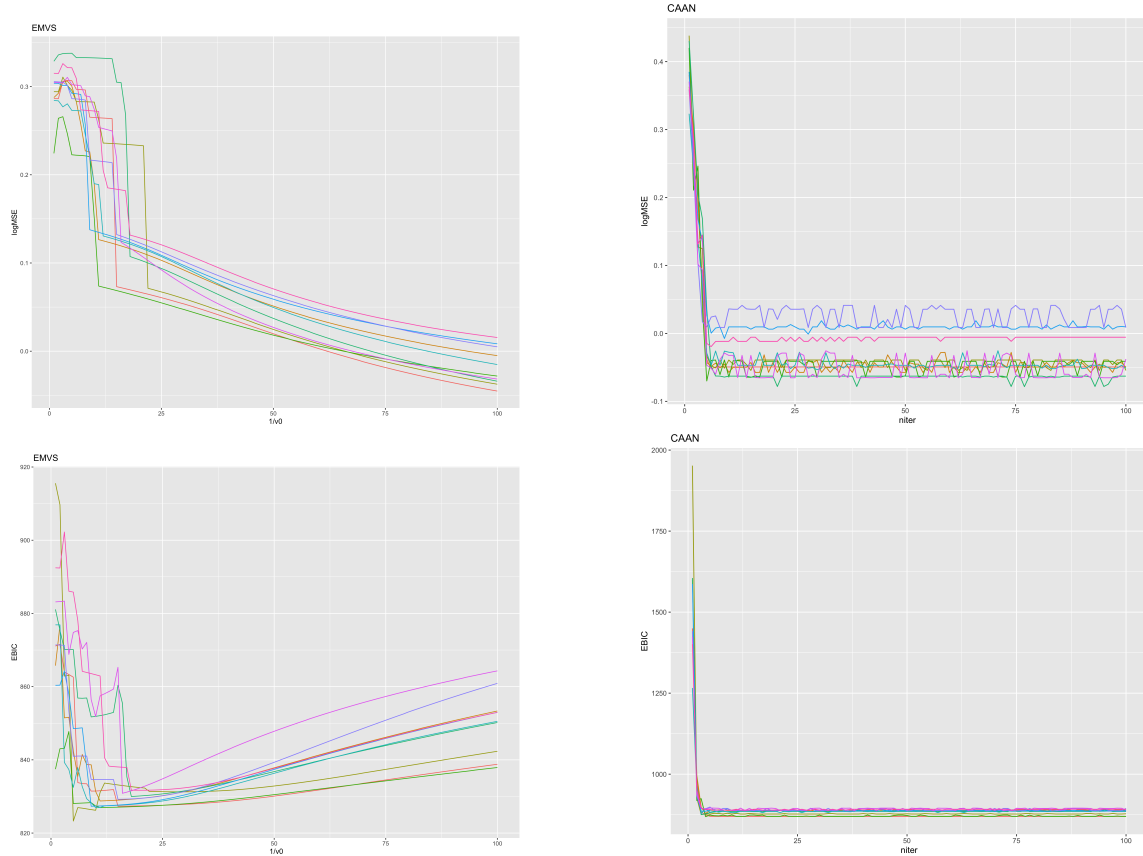


Figure 3: Trace plots of the log-MSE (top row) and EBIC (bottom row) paths from 10 different initial points for the EMVA and the CAAN optimization algorithms, based on a synthetic dataset generated ($n = 120$ and $p = 200$) with the true model 10.

$\nu_1 = 100$ and ν_0^{-1} takes value from $(1, 1000)$. Values of EBIC and log-MSE are evaluated as iteration increases.

Figure 3 shows the optimization paths of EMVS and CAAN. We can see from the results that the optimization paths of EMVS quickly converged to some sub-optimal models, corresponding to different solutions when started with different random initializations. Although both procedures failed to provide consistent results when initialized with different starting configurations, CAAN was more stable than EMVS.

6 Simulation and Real data examples

6.1 Simulation Settings

We generate synthetic data to conduct simulation studies under the linear regression model. We compare the effect of some standard priors including the Bayesian LASSO, the horseshoe and the discrete SpSL with their neuronized counterparts. The optimization-based algorithms EMVS, SSLASSO and CAAN are also included in comparison. We implemented the proposed algorithms in Shin and Liu [2021] and a few R packages are used including Nprior, EMVS and SSLASSO. The notations of the methods for priors in comparison are as follows,

- SpSL-G: Standard discrete SpSL prior, $\pi_0(\theta) = \delta_0(\theta)$, $\pi_1(\theta) = \text{Gaussian}(\theta)$.
- N-SpSL-L(Exact): The neuronized SpSL prior implemented via the exact Gibbs sampler as in Section 5.3.
- N-SpSL-L(RW): The neuronized SpSL prior with a Laplace-like slab via Algorithm 1 using the RWMH to update α .
- N-SpSL-C(RW): The neuronized SpSL prior with a Cauchy slab via Algorithm 1 using the RWMH to update α .
- HS: The Horseshoe prior.
- N-HS(RW): The neuronized Horseshoe prior via Algorithm 1.
- BL: The Bayesian LASSO prior.
- N-BL(RW): The neuronized Bayesian LASSO prior via Algorithm 1.
- N-SpSL(MAP): The neuronized SpSL using the CAAN optimization.
- N-BL(MAP): The neuronized Bayesian LASSO using the CAAN optimization.
- EMVS, SSLASSO : EMVS and SSLASSO.
- LASSO(CV): LASSO with cross validation.

Note that the default τ_w is set to be 1 if not specified. We use $\nu_1 = 10$ for EMVS and $\lambda_1 = 0.1$ for SSLASSO. As is conducted in Shin and Liu [2021], for neuronized priors sampled by MCMC, Bayesian averaging is used for estimation and prediction.

The Median probability model (MPM) is used to do variable selection, which picks variables with estimated posterior inclusion probabilities greater than 0.5, i.e. variables that appear in at least half of the visited models. For neuronized priors, we consider selecting the variable j when

$$1 - \frac{1}{1 + T(\alpha_j - \alpha_0)^2} > 0.5.$$

Shin and Liu [2021] uses a half-collapsed sampling strategy to implement the inference with the SpSL prior to achieve high efficiency. Since no detailed algorithm is provided in the paper, it is hard to figure out and implement in a short time. We use the orthogonal data augmentation method proposed in Ghosh and Clyde [2011] instead to implement models with SpSL prior.

We generated the synthetic data by considering a Toeplitz design. The covariates X_i , $i = 1, \dots, n$ are generated from $N(0, \Sigma)$, where $\Sigma = (\sigma_{lk})$ with $\sigma_{lk} = 0.7^{|l-k|}$ for $l, k = 1, \dots, p$. We generated two scenarios of low-dimensional cases, in which $n = 200$, $p = 50$, and $n = 400$, $p = 100$, respectively. The number of non-zero coefficients is $p/10$ and they take values from ± 0.2 . And another two high-dimensional cases are considered, with $n = 100$, $p = 300$, and $n = 150$, $p = 1000$. First five coefficients are non-zero and they are randomly $\{\pm 0.4, \pm 0.45, \pm 0.5, \pm 0.55, \pm 0.6\}$. The error variance is set to be $\sigma^2 = 1$ for all scenarios.

6.2 Metrics

Shin et al. use three metrics to measure correct sparsity of posterior distributions.

First is mean squared error (MSE) of the predicted coefficient vector from the true data generating parameter vector. MSE for multivariate parameter estimation is calculated as follows.

1. Create estimate $\hat{\beta}$ from sparse prior or equivalent neuronization.
2. Compute the squared Euclidean distance from the estimate to the true data generating parameter vector.
3. Repeat for n dataset instantiations.
4. Calculate the average squared Euclidean distance across all n distances.

Second is cosine similarity. This is calculated in a similar manner as MSE above, except that before step 2 each estimate is also normalized by its own euclidean distance from the origin.

Next is the Mathews correlation coefficient (MCC). This measure is used to determine how correlated each sparse prior is with predicting non-negative coefficients in the true data generating parameter vector. The MCC is defined in terms of tp (true positive predictions), fp (false positive predictions), tn (true negative predictions), and fn (false negative predictions). To clarify, false negatives are instances in which the inference (derived from a particular sparse prior) predicted a zero coefficient for a predictor, but the true coefficient was actually non-zero.

The FP metric, then, is simply the average number of times (over dataset instantiations) a coefficient in the full coefficient vector was estimated to be non-zero, when in fact the true scalar parameter was zero.

The final metric for the simulated dataset was effective sample size (ESS). It is intended to give an indication of the number of uncorrelated MCMC samples. It is calculated as:

$$\frac{N}{1 + 2 * \sum_t^{\infty} \rho(t)}$$

Where N is the number of MCMC samples and $\rho(t)$ is lag- t autocorrelation.

6.3 Simulation results

Tables 1, 2, 3, and 4 show the results of the simulation on the low dimensional and high dimensional settings. In general, the results are identical to the ones in the original paper. We can see from the results that no procedure clearly dominate others in all situations for all criteria. Bayesian averaging results generally have a better performance than the corresponding MAP estimator. And the Lasso-based procedures typically show the best estimation performance under the low-dimensional settings, but they tends to select a large number of false positives. For high dimensional scenarios, the SpSL based priors out performs other methods.

The tables also list the performances of optimization-based SpSL procedures including the CAAN, the EMVS, and the SSLasso. The results show that, overall, the CAAN and the SSLasso significantly outperformed the EMVS algorithms in terms of estimation and model selection.

	MSE	Cos(Angle)	MCC	FP	ESS
SpSL-G	0.228(0.004)	0.573	0.37(0.22)	0.79	3020.3
N-SpSL-L(Exact)	0.228(0.004)	0.579	0.25(0.2)	2.19	2922.7
N-SpSL-C(RW)	0.223(0.004)	0.547	0.23(0.24)	0.17	2985.7
N-HS(RW)	0.223(0.003)	0.585	0.09(0.14)	21.42	2994.4
HS	0.270(0.004)	0.44	0.18(0.07)	3.79	1483.7
BL	0.324(0.013)	0.494	0.212(0.009)	21.290	2348.971
N-BL(RW)	0.328(0.006)	0.586	0.06(0.15)	10.61	2745.4
N-SpSL(MAP)	0.197(0.007)	0.551	0.46(0.2)	2.1	
N-BL(MAP)	0.205(0.008)	0.548	0.43(0.21)	2.49	
EMVS	0.254(0.006)	0.507	0.4(0.18)	5.75	
SSLASSO	0.278(0.01)	0.534	0.39(0.2)	5.84	
Lasso(CV)	0.145(.005)	0.548	0.160(.009)	6.03	

Table 1: Results for low-dimensional setting (n=200, p=50)

6.4 Real data examples

To validate the results in 6.3, Shin et al. applied all procedures to two real-world datasets (described in Section 6.4.1). Since the true data generating parameter is no longer known,

	MSE	Cos(Angle)	MCC	FP	ESS
SpSL-G	0.373(0.004)	0.585	0.45(0.15)	1.03	3018
N-SpSL-L(Exact)	0.362(0.004)	0.591	0.37(0.14)	2.92	2997.8
N-SpSL-C(RW)	0.378(0.004)	0.557	0.31(0.13)	0.18	3012.4
N-HS(RW)	0.361(0.004)	0.63	0.1(0.1)	41.77	2978.2
HS	0.344(0.004)	0.68	0.21(0.14)	2.83	1489.2
BL	0.344(0.010)	0.624	0.253(0.004)	25.470	2355.409
N-BL(RW)	0.472(0.005)	0.691	0.18(0.11)	21.17	2698
N-SpSL(MAP)	0.265(0.008)	0.667	0.61(0.16)	2.64	
N-BL(MAP)	0.266(0.008)	0.672	0.61(0.16)	2.8	
EMVS	0.409(0.006)	0.602	0.54(0.12)	9.04	
SSLASSO	0.35(0.009)	0.656	0.56(0.13)	7.21	
Lasso(CV)	0.238(0.006)	0.645	0.194(0.007)	12.820	

Table 2: Results for low-dimensional setting ($n=400$, $p=100$)

evaluating each procedure was altered slightly from Section 6.2. These changes are described in Section 6.4.2. We then discuss our results as they relate to those reported by Shin et al. in Section 6.4.3

6.4.1 Datasets

The first dataset is the Boston Housing data. This dataset has $n = 506$ datapoints that record information about distinct houses in the Boston area. Shin et al. do not mention which version of this dataset they used, but we obtained it through the `mlbench` package available on CRAN.

This version has 14 features with the *medv* (i.e. median value) feature being the target variable for regression. Shin et al. mention they only used 10 of the other variables for regression, but neglect to mention which ones were used. We therefore used all remaining 13 for our regression experiments on this dataset.

The second dataset used was the Bardet-Beidl syndrome dataset which contains $p = 200$ features indicating gene expression information in $n = 120$ rats. Because this dataset has $p > n$ features, inference with most of the tested algorithms took longer than with the Boston Housing dataset.

6.4.2 Evaluation

Shin et al. used three metrics to evaluate shrinkage prior performance on the datasets.

The first is mean-squared *predictive* error (MSPE). This is similar to the MSE described above, but the squared error is taken with respect to the regression variable rather than the full coefficient vector. As such, the MSPE takes it's average over every predicted datapoint

	MSE	Cos(Angle)	MCC	FP	ESS
SpSL-G	1.468(0.061)	0.675	0.31(0.15)	13.97	2803.3
N-SpSL-L(Exact)	1.182(0.021)	0.625	0.24(0.16)	6.7	2996.6
N-SpSL-C(RW)	1.093(0.021)	0.63	0.3(0.22)	0.55	2875.1
N-HS(RW)	1.387(0.023)	0.634	0.05(0.05)	143.59	2954.8
HS	1.66(0.034)	0.71	0.54(0.01)	99.35	2563.2
BL	0.931(0.020)	0.586	0.073(0.003)	4.210	2395.840
N-BL(RW)	1.721(0.021)	0.527	0.06(0.06)	90.55	2860.3
N-SpSL(MAP)	1.43(0.047)	0.475	0.27(0.13)	13.35	
N-BL(MAP)	1.436(0.049)	0.482	0.26(0.13)	13.28	
EMVS	1.588(0.02)	0.318	0.28(0.1)	21.13	
SSLASSO	0.994(0.033)	0.526	0.4(0.21)	8.14	
Lasso(CV)	0.896(0.022)	0.579	0.072(0.004)	14.200	

Table 3: Results for high-dimensional setting (n=100, p=300)

in the dataset. This is in contrast to MSE which only measured the squared error between the one parameter vector estimate obtained from the posterior and the true vector. Since Shin et al. repeated their experiments 100 times, each MSPE estimate was averaged again.

Shin et al. also used cosine angle as a metric in their real-world dataset experiments. Similar to the MSE measure, the cosine similarity could no longer be made with respect to the coefficient vector. If \mathbf{y} is the dependent regression variable in the dataset and the $\hat{\mathbf{y}}$ is the prediction, then the cosine similarity used here in $\mathbf{y}^T \hat{\mathbf{y}} / (\|\mathbf{y}\| \cdot \|\hat{\mathbf{y}}\|)$. This measure indicates the correlation between the prediction and actual response variable

The final metric used is *model size*. This is simply the average number of non-zero coefficients chosen by each procedure.

6.4.3 Discussion

The results for the experiments on the Boston Housing dataset can be found in Table 5 and the Bardet-Biedl experiments in Table 6.

In the Boston housing experiments there is a striking similarity in all neuronized procedures. We hypothesise that this is due mainly to the similar optimization method that is used in each of them. It is also noteworthy that the model sizes of each neuronized implementation are substantially less than both the Bayesian Lasso (BL) and Lasso procedures.

In the Bardet-Biedl experiments, however, there was a lot more variability in the model sizes induced by each procedure. This trend is matched in the MSPE for each procedure, which saw scores as low as .017 for the neuronized horseshoe prior, all the way to .907 for the Bayesian Lasso. The Bayesian Lasso, in particular, appeared to struggle with finding good regression coefficient estimates as evidenced by its cosine angle of .089. Further study is needed to explain the behavior of the Bayesian Lasso in this setting.

	MSE	Cos(Angle)	MCC	FP	ESS
SpSL-G	2.921(0.154)	0.538	0.13(0.13)	245.52	223.3
N-SpSL-L(Exact)	0.983(0.026)	0.659	0.27(0.14)	8.23	156.1
N-SpSL-C(RW)	0.956(0.026)	0.624	0.35(0.21)	0.54	152.6
N-HS(RW)	1.973(0.031)	0.656	0.03(0.03)	482.06	132.4
HS	0.259(0.032)	0.594	0.12(0.16)	60.89	2678.1
BL	1.163(0.032)	0.378	0.238(0.016)	13.610	146.434
N-BL(RW)	1.828(0.023)	0.51	0.13(0.03)	312.13	132.4
N-SpSL(MAP)	0.239(0.087)	0.53	0.33(0.16)	9.77	
N-BL(MAP)	0.272(0.094)	0.494	0.29(0.13)	12.82	
EMVS	1.32(0.01)	0.242	0.74(0.17)	0.08	
SSLASSO	1.063(0.037)	0.496	0.36(0.19)	12.2	
Lasso(CV)	0.869(0.023)	0.626	0.375(0.019)	21.350	

Table 4: Results for high-dimensional setting (n=150, p=1000)

7 Conclusion

By replicating the experiments in Shin and Liu’s essayShin and Liu [2021], it is found that there are still some differences between our experimental results for simulated and real data and the results given by the authors. However, while the same analysis program was used, the difference between the results for the real data is relatively small, while the difference between the results for the simulated data is large, which could indicate that the difference in the results may come from the difference between the simulated sample data set and the original work, while the data processing process itself is not too problematic.

From the results in the table, we can see that the neuronized prior usually has a smaller ESS than the original prior, which proves that the neuronizing operation can improve the efficiency of analyzing the data, and the difference is most obvious in horseshoe prior.

As in the Shin and Liu’s essay, the MSE of lasso shows the best estimation performance, but the false positive is higher. By comparing MCC, we found that SpSL prior generally has a better model selection process in high dimensional settings, while horseshoe and its neuronized prior has a preferred result in low-dimensional setting.

8 Chris Crabtree (My Own) - Contribution

In terms of experiments, I helped by implementing the Lasso and Baysian Lasso procedures. I also started on the SSLASSO procedure, but ran out of time. Additionally, I took charge of navigating the evaluation produres for the experiments. The original authors were quite terse in their descriptions and left several experimental settings ambiguous or underspecified. I was tasked, then, with exploring several possible options for evaluation

	MSPE	Cos(Angle)	MS
SpSL-G	24.577(0.951)	0.843	5.17
N-SpSL-L(Exact)	24.246(0.995)	0.846	8.52
N-SpSL-C(RW)	24.397(1.005)	0.846	5.24
N-HS(RW)	24.493(0.962)	0.844	7.74
HS	24.377(1.23)	0.843	7.764
BL	26.7(0.9)	0.857	12.260
N-BL(RW)	24.377(1.03)	0.846	7.03
N-SpSL(MAP)	23.977(0.987)	0.845	7.68
N-BL(MAP)	23.974(0.986)	0.845	8.44
Lasso(CV)	26.7(0.9)	0.857	12.260

Table 5: Results for Boston housing dataset

	MSPE	Cos(Angle)	MS
SpSL-G	0.018(0.001)	0.572	99.38
N-SpSL-L(Exact)	0.049(0.003)	0.674	110.12
N-SpSL-C(RW)	0.026(0.002)	0.645	1.81
N-HS(RW)	0.017(0.001)	0.698	98.13
HS	0.018(0.008)	0.689	87.63
BL	0.907(0.072)	0.083	30.080
N-BL(RW)	0.022(0.001)	0.649	59.19
N-SpSL(MAP)	0.735(0.064)	0.589	1.93
N-BL(MAP)	0.844(0.253)	0.601	62.06
Lasso(CV)	0.454(0.026)	0.658	32.010

Table 6: Results for Bardet-Biedl dataset

to see which type produced the results seen in the paper.

Since I made those contributions, I naturally was responsible for the bulk of the writing of the report w.r.t. the lasso procedures and the presentation of experimental results. I also helped create the introductory slides for the video presentation.

9 Yunshan Duan - Contribution

Yunshan was for sure the MVP of the group. She implemented the majority of the algorithms and also wrote the largest chunks of the paper

10 Lanxin Yang - Contribution

Lanxin helped by implementing the Horseshoe prior and filling out large portions of the Standard Prior section of the report. She was also responsible for the video editing.

References

- Anirban Bhattacharya, Antik Chakraborty, and Bani K Mallick. Fast sampling with gaussian scale mixture priors in high-dimensional regression. *Biometrika*, page asw042, 2016.
- Carlos M Carvalho, Nicholas G Polson, and James G Scott. The horseshoe estimator for sparse signals. *Biometrika*, 97(2):465–480, 2010.
- Ismaël Castillo and Aad van der Vaart. Needles and straw in a haystack: Posterior concentration for possibly sparse sequences. *The Annals of Statistics*, 40(4):2069–2101, 2012.
- Ismaël Castillo, Johannes Schmidt-Hieber, and Aad Van der Vaart. Bayesian linear regression with sparse priors. *The Annals of Statistics*, 43(5):1986–2018, 2015.
- Xiaohui Chen, Z Jane Wang, and Martin J McKeown. A bayesian lasso via reversible-jump mcmc. *Signal Processing*, 91(8):1920–1932, 2011.
- Edward I George and Robert E McCulloch. Variable selection via gibbs sampling. *Journal of the American Statistical Association*, 88(423):881–889, 1993.
- Joyee Ghosh and Merlise A Clyde. Rao–blackwellization for bayesian variable selection and model averaging in linear and binary regression: A novel data augmentation approach. *Journal of the American Statistical Association*, 106(495):1041–1052, 2011.
- Peter D Hoff. Lasso, fractional norm and structured sparse estimation using a hadamard product parametrization. *Computational Statistics & Data Analysis*, 115:186–198, 2017.
- S. Kirkpatrick, C. D. Gelatt, and M. P. Vecchi. Optimization by simulated annealing. *Science*, 220(4598):671–680, 1983. doi: 10.1126/science.220.4598.671.
- Jun S Liu and Chiara Sabatti. Generalised gibbs sampler and multigrid monte carlo for bayesian computation. *Biometrika*, 87(2):353–369, 2000.
- K. L. Mengersen and R. L. Tweedie. Rates of convergence of the hastings and metropolis algorithms. *Journal of the Royal Statistical Society: Series B (Methodological)*, 58(1):101–121, 1996.
- Toby J Mitchell and John J Beauchamp. Bayesian variable selection in linear regression. *Journal of the american statistical association*, 83(404):1023–1032, 1988.
- Trevor Park and George Casella. The bayesian lasso. *Journal of the American Statistical Association*, 103(482):681–686, 2008.
- Veronika Ročková and Edward I George. Emvs: The em approach to bayesian variable selection. *Journal of the American Statistical Association*, 109(506):828–846, 2014.

- Veronika Ročková and Edward I George. The spike-and-slab lasso. *Journal of the American Statistical Association*, 113(521):431–444, 2018.
- James G Scott and James O Berger. Bayes and empirical-bayes multiplicity adjustment in the variable-selection problem. *The Annals of Statistics*, pages 2587–2619, 2010.
- Minsuk Shin and Jun S Liu. Neuronized priors for bayesian sparse linear regression. *Journal of the American Statistical Association*, (just-accepted):1–43, 2021.
- Robert Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B (Methodological)*, 58(1):267–288, 1996.